

Group: Zeru Zhou's Group

Name: Zeru Zhou

Email: zeruzhou9@gmail.com

Country: United States

University: University of Southern California

Specialization: Data Science

Github Repo:

Problem Description

In this project, I need to build binary classification machine learning models to predict if the bank clients will renew the term deposit or not in order to make corresponding strategies to maintain clients.

Dataset Information

The dataset contains 3 parts: Client information, campaign information and social/economical context.

There are total 41188 instances and 20 columns.

The target response is whether the bank client will renew the term deposit.

Library importing

```
In [1]: import pandas as pd
import numpy as np
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.svm import SVC
from sklearn.svm import LinearSVC
from sklearn.naive_bayes import GaussianNB, MultinomialNB
from sklearn.linear_model import RidgeClassifierCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from imblearn.pipeline import Pipeline
#from sklearn.pipeline import Pipeline
from imblearn.over_sampling import SMOTE
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, precision_score,
```

```

from sklearn.model_selection import train_test_split
import warnings
from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings("ignore", category=ConvergenceWarning)
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import MinMaxScaler
from sklearn.feature_selection import RFECV

```

Data uploading

```
In [2]: my_df = pd.read_csv('../data/bank-additional/bank-additional-full.csv', sep=';')
```

```
In [3]: my_df.shape
```

```
Out[3]: (41188, 21)
```

```
In [4]: my_df.head()
```

```
Out[4]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	cam
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	
1	57	services	married	high.school	unknown		no	no	telephone	may	mon	...
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	

5 rows x 21 columns

```
In [5]: my_df['y'].value_counts()
```

```
Out[5]: no      36548
yes       4640
Name: y, dtype: int64
```

```
In [6]: my_df.isnull().any()
```

```
Out[6]: age                False
job                False
marital            False
education          False
default            False
housing            False
loan               False
contact            False
month              False
day_of_week        False
duration           False
campaign           False
pdays            False
previous           False
poutcome           False
emp.var.rate       False
cons.price.idx     False
cons.conf.idx      False
euribor3m          False
nr.employed        False
y                  False
dtype: bool
```

```
In [7]: sub = []
        for i in my_df.isnull().any().keys():
            if my_df.isnull().any()[i] == False:
                sub.append(i)
        sub
```

```
Out[7]: ['age',
         'job',
         'marital',
         'education',
         'default',
         'housing',
         'loan',
         'contact',
         'month',
         'day_of_week',
         'duration',
         'campaign',
         'pdays',
         'previous',
         'poutcome',
         'emp.var.rate',
         'cons.price.idx',
         'cons.conf.idx',
         'euribor3m',
         'nr.employed',
         'y']
```

As we can see, this is an imbalanced dataset and it has NA values. We need to clean the dataset before moving to any analysis and model building.

Data Cleaning and imputation

Drop duplicated rows

```
In [8]: my_df.duplicated(subset= sub).value_counts()
```

```
Out[8]: False    41176
        True      12
        dtype: int64
```

```
In [9]: my_df = my_df.drop_duplicates()
        my_df.shape
```

```
Out[9]: (41176, 21)
```

Impute missing values

First, find categorical variables

```
In [78]: my_df.dtypes
```

```
Out[78]: age                int64
         job                object
         marital            object
         education          object
         default            object
         balance            float64
```

```

housing      object
loan         object
contact      object
day          float64
month        object
duration     int64
campaign     int64
pdays       int64
previous     int64
poutcome     object
y            object
day_of_week  object
emp.var.rate float64
cons.price.idx float64
cons.conf.idx float64
euribor3m    float64
nr.employed  float64
dtype: object

```

```

In [80]: my_dict = {}
         for n,i in enumerate(my_df.columns):
             my_dict[i] = my_df.dtypes[n]
         type_df = pd.DataFrame(my_dict, index=['dtype'])
         type_df

```

```

Out[80]:
      age  job  marital  education  default  balance  housing  loan  contact  day  month  durati
dtype  int64  object  object    object  object  float64    object  object  object  float64  object  int

```

```

In [86]: type_df = type_df.T
         type_df

```

```

Out[86]:
      dtype
age      int64
job      object
marital  object
education  object
default  object
balance  float64
housing  object
loan     object
contact  object
day      float64
month    object
duration int64
campaign int64
pdays   int64
previous int64
poutcome object
y        object
day_of_week  object

```

```
emp.var.rate    float64
cons.price.idx  float64
cons.conf.idx   float64
euribor3m       float64
nr.employed     float64
```

```
In [90]: col = type_df.loc[type_df['dtype'] == 'object']
categorical = list(col.index)
categorical.remove('y')
categorical
```

```
Out[90]: ['job',
'marital',
'education',
'default',
'housing',
'loan',
'contact',
'month',
'poutcome',
'day_of_week']
```

```
In [91]: df = pd.get_dummies(my_df, columns=categorical)
```

```
In [93]: df.head()
```

Out[93]:

	age	balance	day	duration	campaign	pdays	previous	y	emp.var.rate	cons.price.idx	cons.conf.id
0	58	2143.0	5.0	261	1	-1	0	no	NaN	NaN	NaN
1	44	29.0	5.0	151	1	-1	0	no	NaN	NaN	NaN
2	33	2.0	5.0	76	1	-1	0	no	NaN	NaN	NaN
3	47	1506.0	5.0	92	1	-1	0	no	NaN	NaN	NaN
4	33	1.0	5.0	198	1	-1	0	no	NaN	NaN	NaN

```
In [94]: x,y = df.drop(columns=['y'], df['y'])
```

First, we can try simple imputer like filling the value with mean

```
In [95]: x
```

Out[95]:

	age	balance	day	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.ic
0	58	2143.0	5.0	261	1	-1	0	NaN	NaN	NaN
1	44	29.0	5.0	151	1	-1	0	NaN	NaN	NaN
2	33	2.0	5.0	76	1	-1	0	NaN	NaN	NaN
3	47	1506.0	5.0	92	1	-1	0	NaN	NaN	NaN
4	33	1.0	5.0	198	1	-1	0	NaN	NaN	NaN
...
41183	73	NaN	NaN	334	1	999	0	-1.1	94.767	-50
41184	46	NaN	NaN	383	1	999	0	-1.1	94.767	-50

41185	56	NaN	NaN	189	2	999	0	-1.1	94.767	-50
41186	44	NaN	NaN	442	1	999	0	-1.1	94.767	-50
41187	74	NaN	NaN	239	3	999	1	-1.1	94.767	-50

86387 rows × 73 columns

```
In [104... x = x.drop(columns=['duration'])
```

```
In [105... x_simple = SimpleImputer(strategy='mean').fit_transform(x)
pd.DataFrame(x_simple).head()
```

Out[105]:	0	1	2	3	4	5	6	7	8	9	10	11	12	13
0	58.0	2143.0	5.0	1.0	-1.0	0.0	0.081922	93.57572	-40.502863	3.621293	5167.03487	0.0	0.0	0.0
1	44.0	29.0	5.0	1.0	-1.0	0.0	0.081922	93.57572	-40.502863	3.621293	5167.03487	0.0	0.0	0.0
2	33.0	2.0	5.0	1.0	-1.0	0.0	0.081922	93.57572	-40.502863	3.621293	5167.03487	0.0	0.0	1.0
3	47.0	1506.0	5.0	1.0	-1.0	0.0	0.081922	93.57572	-40.502863	3.621293	5167.03487	0.0	1.0	0.0
4	33.0	1.0	5.0	1.0	-1.0	0.0	0.081922	93.57572	-40.502863	3.621293	5167.03487	0.0	0.0	0.0

Since this is high-dimensional dataset with almost 100k instances, using simple imputer may lead to huge error. So we try iterative imputer using round robin algorithm

```
In [134... x_iter = IterativeImputer(n_nearest_features=50).fit_transform(x)
x_iter_df = pd.DataFrame(x_iter)
```

```
In [108... x_iter_df.head()
```

Out[108]:	0	1	2	3	4	5	6	7	8	9	10	11	12
0	58.0	2143.0	5.0	1.0	-1.0	0.0	-0.001759	93.460615	-37.686446	3.690600	5166.507093	0.0	0.0
1	44.0	29.0	5.0	1.0	-1.0	0.0	0.013415	93.471578	-38.045953	3.665590	5165.584233	0.0	0.0
2	33.0	2.0	5.0	1.0	-1.0	0.0	-0.012448	93.451250	-38.446562	3.637786	5166.385384	0.0	0.0
3	47.0	1506.0	5.0	1.0	-1.0	0.0	0.008979	93.502296	-38.150717	3.642176	5163.373385	0.0	1.0
4	33.0	1.0	5.0	1.0	-1.0	0.0	-0.024682	93.495518	-37.526338	3.675858	5162.429888	0.0	0.0

After filling all the NAs, the next step is to eliminate some outliers

```
In [109... x_iter_df.describe()
```

Out[109]:	0	1	2	3	4	5
count	86387.000000	86387.000000	86387.000000	86387.000000	86387.000000	86387.000000
mean	40.501314	1068.949277	14.197018	2.670437	479.792504	0.386181
std	10.534612	2267.890700	6.665170	2.947981	483.824356	1.131173
min	17.000000	-8019.000000	1.000000	1.000000	-1.000000	0.000000
25%	32.000000	163.000000	10.445861	1.000000	-1.000000	0.000000
50%	39.000000	584.000000	13.730860	2.000000	246.000000	0.000000

75%	48.000000	1211.586290	17.166562	3.000000	999.000000	0.000000	1.1
max	98.000000	102127.000000	36.740440	63.000000	999.000000	275.000000	10.8

In [126... y

Out[126]:

0	no
1	no
2	no
3	no
4	no
...	
41183	yes
41184	no
41185	no
41186	yes
41187	no

Name: y, Length: 86387, dtype: object

In [135... x_iter_df['y'] = y.values

In [17]:

```

my_list = []
for key, value in my_df.dtypes.items():
    if value == 'object':
        my_list.append(key)
my_list.remove('y')
my_list

```

Out[17]:

```

['job',
 'marital',
 'education',
 'default',
 'housing',
 'loan',
 'contact',
 'month',
 'day_of_week',
 'poutcome']

```

In [18]:

```

my_df = pd.get_dummies(my_df, columns=my_list)
my_df

```

Out[18]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m
0	56	261	1	999	0	1.1	93.994	-36.4	4.857
1	57	149	1	999	0	1.1	93.994	-36.4	4.857
2	37	226	1	999	0	1.1	93.994	-36.4	4.857
3	40	151	1	999	0	1.1	93.994	-36.4	4.857
4	56	307	1	999	0	1.1	93.994	-36.4	4.857
...
41181	37	281	1	999	0	-1.1	94.767	-50.8	1.028
41182	29	112	1	9	1	-1.1	94.767	-50.8	1.028
41184	46	383	1	999	0	-1.1	94.767	-50.8	1.028
41185	56	189	2	999	0	-1.1	94.767	-50.8	1.028
41186	44	442	1	999	0	-1.1	94.767	-50.8	1.028

40428 rows x 64 columns

```
In [19]: for i in my_df.columns:
        if i == 'y':
            continue
        q_low = my_df[i].quantile(0.01)
        q_high = my_df[i].quantile(0.99)
        my_df = my_df.loc[(my_df[i] <= q_high) & (my_df[i] >= q_low)]
```

```
In [20]: my_df.shape
```

```
Out[20]: (36103, 64)
```

```
In [30]: my_df = my_df.drop(columns=['duration'])
        my_df.shape
```

```
Out[30]: (36103, 63)
```

Here above, For each feature, I only dropped very extreme values that not belong to the central 99%. It is also feasible to use q1, q3, and 1.5 IQR to detect outliers but I don't want that much data lose.

Model Building

For each model, we'll build pipeline for feature selection and hyperparameter tuning, using cross validation. Also SMOTE is included since this is imbalanced dataset.

Split the dataset into train/test

```
In [31]: x, y = my_df.drop(columns=['y']), my_df['y']
```

```
In [32]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.1, random_state=42)
        x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
Out[32]: ((32492, 62), (3611, 62), (32492,), (3611,))
```

Logistic Regression

For Logistic Regression, we can perform feature selection by using l1 penalization, it is a shrink method by reducing the coefficients of irrelevant variables to 0. We can also use other methods but this is the most convenient way. We do feature selection together with hyperparameter tuning.

```
In [33]: log = LogisticRegression(penalty='l1', solver='saga', max_iter=300)
```

```
In [34]: scale = MinMaxScaler()
```

```
In [35]: imbalance = SMOTE()
```

```
In [36]: param = {'model__C': np.logspace(-3, 3, num=50)}
```



```
In [37]: x_train_scaled = scale.fit_transform(x_train)
```

```
In [38]: pipe = Pipeline(steps=[('smote', imbalance), ('model', log)])
```

```
In [39]: clf_log = GridSearchCV(pipe, param, n_jobs=-1).fit(x_train_scaled, y_train)
```

[illegible]

```
In [40]: x_test_scaled = scale.fit_transform(x_test)
```

```
In [41]: y_pred = clf_log.predict(x_test_scaled)
```

```
In [42]: acc = accuracy_score(y_test, y_pred)
print(f'The accuracy score is {acc}')
```

The accuracy score is 0.7864857380227084

```
In [43]: pre = precision_score(y_test, y_pred, pos_label='no')
print(f'The precision score is {pre}')
```

The precision score is 0.9589578872234118

```
In [44]: rec = recall_score(y_test, y_pred, pos_label='no')
print(f'The recall score is {rec}')
```

The recall score is 0.8037690696978762

```
In [45]: cm = confusion_matrix(y_test, y_pred)
          print(cm)
```

```
[[2687  656]
 [ 115  153]]
```

Try RFE instead of l1-penalty

```
In [46]: log = LogisticRegression(max_iter=5000, penalty='none')
```

```
In [47]: feature, target = imbalance.fit_resample(x_train, y_train)

In [48]: feature_scaled = scale.fit_transform(feature)

In [49]: clf_log_rfe = RFECV(log, cv=5, n_jobs=-1).fit(feature_scaled, target)

In [50]: y_pred = clf_log_rfe.predict(x_test_scaled)
```

```
In [51]: acc = accuracy_score(y_test, y_pred)
print(f'The accuracy score is {acc}')
```

The accuracy score is 0.9277208529493215

```
In [52]: pre = precision_score(y_test, y_pred, pos_label='no')
print(f'The precision score is {pre}')
```

The precision score is 0.9331084879145587

```
In [53]: rec = recall_score(y_test, y_pred, pos_label='no')
print(f'The recall score is {rec}')
```

The recall score is 0.9931199521387974

```
In [54]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[3320   23]
 [ 238   30]]
```

Without class imbalance treatment

```
In [55]: log = LogisticRegression(max_iter=5000)
```

```
In [56]: param = {'C': np.logspace(-3, 3, num=50)}
```

```
In [57]: clf_log_1 = GridSearchCV(log, param, n_jobs=-1).fit(x_train_scaled, y_train)
```

```
/Users/zhouzeru/opt/anaconda3/lib/python3.9/site-packages/joblib/externals/loky/process_executor.py:702: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.
  warnings.warn(
```

```
In [58]: y_pred = clf_log_1.predict(x_test_scaled)
```

```
In [59]: acc = accuracy_score(y_test, y_pred)
print(f'The accuracy score is {acc}')
```

The accuracy score is 0.9304901689282747

```
In [60]: pre = precision_score(y_test, y_pred, pos_label='no')
print(f'The precision score is {pre}')
```

The precision score is 0.9342696629213483

```
In [61]: rec = recall_score(y_test, y_pred, pos_label='no')
print(f'The recall score is {rec}')
```

The recall score is 0.9949147472330242

```
In [62]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[3326   17]
 [ 234   34]]
```

K-Neighbor

```
In [63]: knn = KNeighborsClassifier(weights='uniform')
```

```
In [64]: params = {'n_neighbors':[int(x) for x in np.linspace(1,20,num=20)]}
```

```
In [65]: clf_knn = GridSearchCV(knn, params, n_jobs=-1).fit(feature_scaled, target)
```

```
In [66]: y_pred = clf_knn.predict(x_test_scaled)
```

```
In [67]: acc = accuracy_score(y_test, y_pred)
print(f'The accuracy score is {acc}')
```

The accuracy score is 0.910828025477707

```
In [68]: pre = precision_score(y_test, y_pred, pos_label='no')
print(f'The precision score is {pre}')
```

The precision score is 0.934176487496407

```
In [69]: rec = recall_score(y_test, y_pred, pos_label='no')
print(f'The recall score is {rec}')
```

The recall score is 0.9721806760394855

```
In [70]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[3250   93]
 [ 229   39]]
```

Ensemble Tree

```
In [71]: rdf = RandomForestClassifier()
```

```
In [72]: params = {'ccp_alpha': np.logspace(-3, 3, num=20)}
```

```
In [73]: clf_rdf = GridSearchCV(rdf, params, n_jobs=-1).fit(feature_scaled, target)
```

```
In [74]: y_pred = clf_rdf.predict(x_test_scaled)
```

```
In [75]: acc = accuracy_score(y_test, y_pred)
print(f'The accuracy score is {acc}')
```

The accuracy score is 0.8759346441428967

```
In [76]: pre = precision_score(y_test, y_pred, pos_label='no')
print(f'The precision score is {pre}')
```

The precision score is 0.9570571518787496

```
In [77]: rec = recall_score(y_test, y_pred, pos_label='no')
print(f'The recall score is {rec}')
```

The recall score is 0.9066706551002094

```
In [78]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[3031  312]
 [ 136  132]]
```

Boosting Tree

```
In [79]: xgb = XGBClassifier(eta = 0.01, objective = 'binary:logistic')
```

```
In [80]: params = {'reg_alpha': np.logspace(-3, 3, 10)}
```

```
In [81]: def change(num):  
    if num == 'no':  
        return 0  
    else:  
        return 1
```

```
In [82]: target_num = [change(i) for i in target]
```

```
In [83]: clf_xgb = GridSearchCV(xgb, params, n_jobs=-1).fit(feature_scaled, target_num)
```

```
In [84]: y_pred = clf_xgb.predict(x_test_scaled)
```

```
In [85]: y_test_num = [change(i) for i in y_test]
```

```
In [86]: acc = accuracy_score(y_test_num, y_pred)  
print(f'The accuracy score is {acc}')
```

The accuracy score is 0.8726114649681529

```
In [87]: pre = precision_score(y_test_num, y_pred)  
print(f'The precision score is {pre}')
```

The precision score is 0.2818181818181818

```
In [88]: rec = recall_score(y_test_num, y_pred)  
print(f'The recall score is {rec}')
```

The recall score is 0.4626865671641791

```
In [89]: cm = confusion_matrix(y_test_num, y_pred)  
print(cm)
```

```
[[3027  316]  
 [ 144  124]]
```

SVM

```
In [92]: svc = SVC()
```

```
In [98]: params = {'C': np.logspace(-3, 2, 10)}
```

```
In [99]: clf_svm = GridSearchCV(svc, params, n_jobs=-1).fit(feature_scaled, target)
```

```
/Users/zhouzeru/opt/anaconda3/lib/python3.9/site-packages/joblib/externals/loky/process_executor.py:702: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.  
warnings.warn()
```

```
In [100]: y_pred = clf_svm.predict(x_test_scaled)
```

```
In [101]: acc = accuracy_score(y_test, y_pred)  
print(f'The accuracy score is {acc}')
```

The accuracy score is 0.9293824425366934

```
In [102... pre = precision_score(y_test, y_pred, pos_label='no')
print(f'The precision score is {pre}')
```

The precision score is 0.9356659142212189

```
In [103... rec = recall_score(y_test, y_pred, pos_label='no')
print(f'The recall score is {rec}')
```

The recall score is 0.9919234220759796

```
In [104... cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[3316  27]
 [ 228  40]]
```

As results above, I've trialed several ML models with feature engineering & hyperparameter tuning. Logistic regression had the best performance in contrast to SVM, Bayesian model, and ensemble tree models.

According to the models above and the metrics we used, logistic regression with recursive feature elimination and support vector machine with RBF kernel had the greatest accuracy while considerable precision/recall score. When converting these ML metrics into business metrics, this filtering model that select potential clients could significantly improve our targeting efficiency. Reducing the number of targeted clients while still convincing a lot of them to subscribe, the targeting efficiency could be boosted.